

RUN1 native machinery · observed events, not causal attributions

A Recorded mechanism use

Native islands

[exercised](#)

4 configured; 0 migration events

Parent inspirations

[exercised](#)

Archive contexts 1; top-k 0

Mutation operators

[exercised](#)

diff: 1, full: 1; crossover not exercised

Model bandit

[two arms queried with outcome feedback](#)

2 mutation arms queried; cost coefficient 0

Native novelty

[discovery judge exercised](#)

3 embedding calls; 2 gate decisions; 0 rejected

Meta recommendations

[recommendations consumed](#)

5 calls; 1 contexts consumed advice

Prompt coevolution

[unexercised before stop](#)

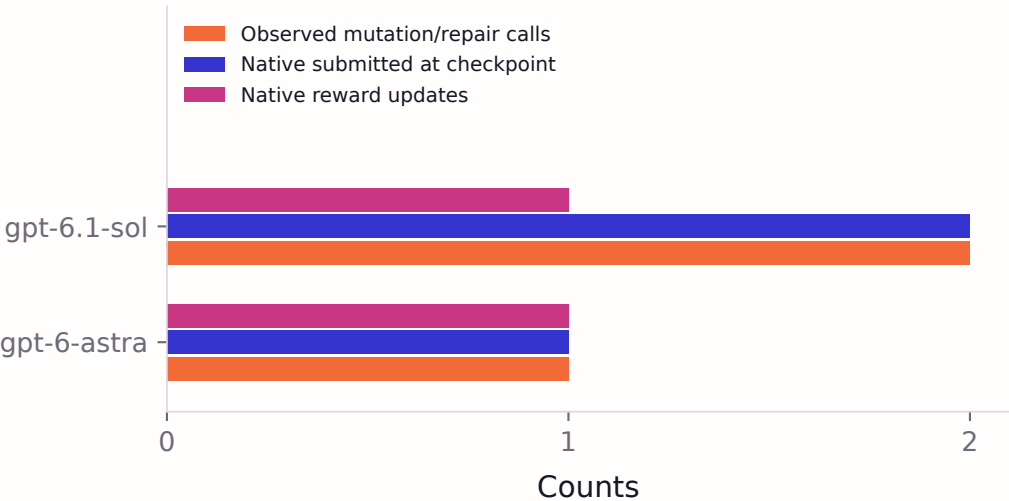
1 stored prompts; 2 program credits

Checkpoint / resume

[actually restored](#)

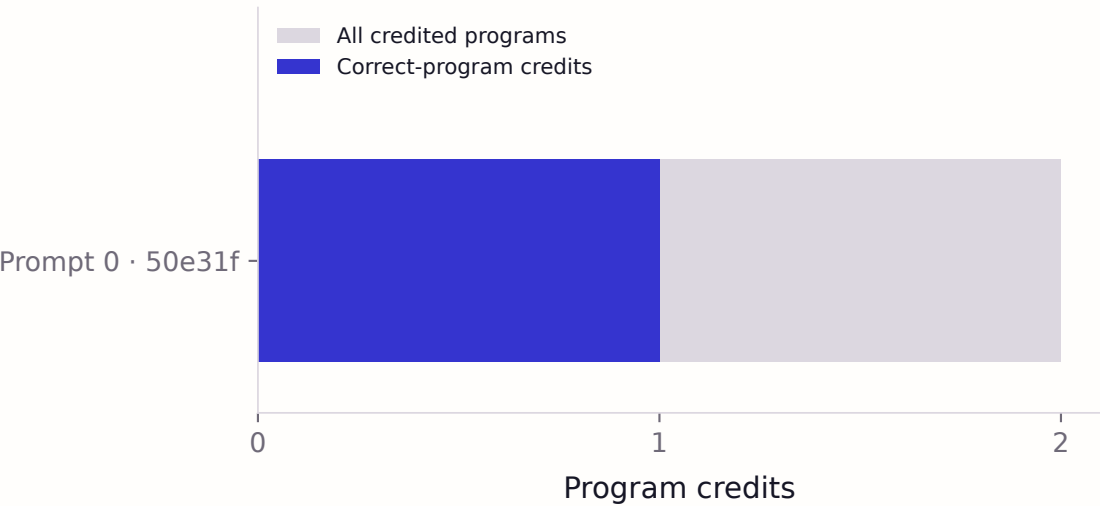
1 held slots without source; 0 remote repair calls

B Arm allocation and reward updates



Checkpoint boundary: launcher_finally
An active call can postdate the saved UCB state.

C Native prompt credit



Native prompt fitness averages correct-program percentiles.
Counts include failures; this is not failure-inclusive task fitness.

Readiness and native novelty fixture calls are separate from discovery; a fixture does not prove the discovery judge ran.
Local BGE embeddings truncate to 10,000 source characters and are not validated semantic code novelty.
Native UCB reward updates are shifted/scaled feedback; allocation is not a controlled comparison of model quality.
Pending means configured but not exercised at this snapshot. One search cannot isolate a mechanism’s benefit or reliable discovery.