

RUN1 native machinery · observed events, not causal attributions

A Recorded mechanism use

Native islands

exercised

4 configured; 0 migration events

Parent inspirations

pending

Archive contexts 0; top-k 0

Mutation operators

exercised

diff: 1, full: 1

Model bandit

two arms queried with outcome feedback

2 mutation arms queried; cost coefficient 0

Native novelty

discovery judge exercised

3 embedding calls; 2 gate decisions; 0 rejected

Meta recommendations

calls exercised

5 calls; 0 contexts consumed advice

Prompt coevolution

pending

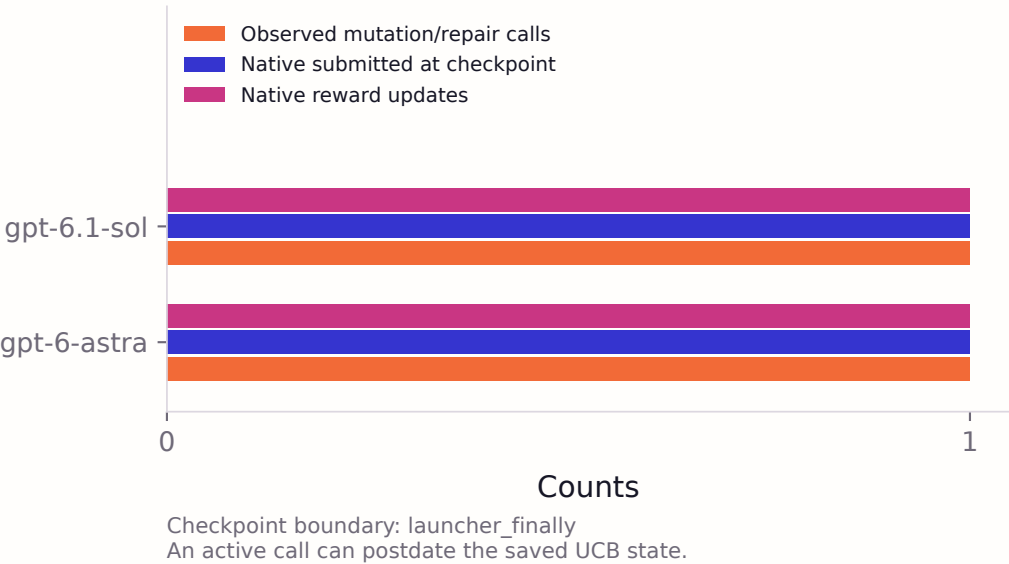
1 stored prompts; 2 program credits

Checkpoint / resume

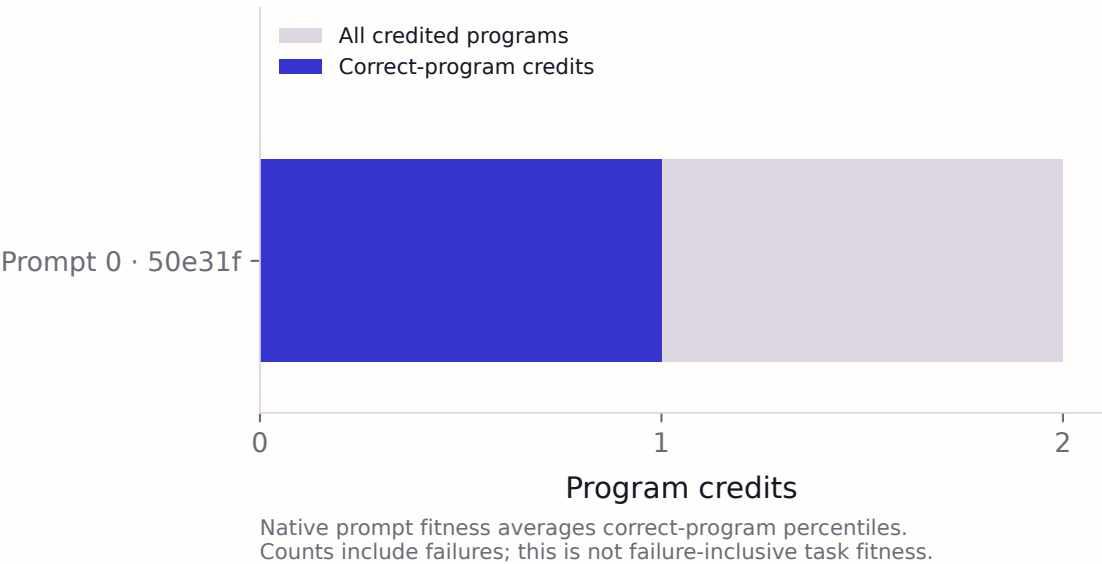
checkpoint available not live resumed

0 held proposals; 0 patch repair attempts

B Arm allocation and reward updates



C Native prompt credit



Readiness and native novelty fixture calls are separate from discovery; a fixture does not prove the discovery judge ran.
Local BGE embeddings truncate to 10,000 source characters and are not validated semantic code novelty.
Native UCB reward updates are shifted/scaled feedback; allocation is not a controlled comparison of model quality.
All mechanisms are enabled together. One search cannot establish reliable discovery or isolate a mechanism’s benefit.