



RUN1 resources · completed usage and local CPU are different clocks

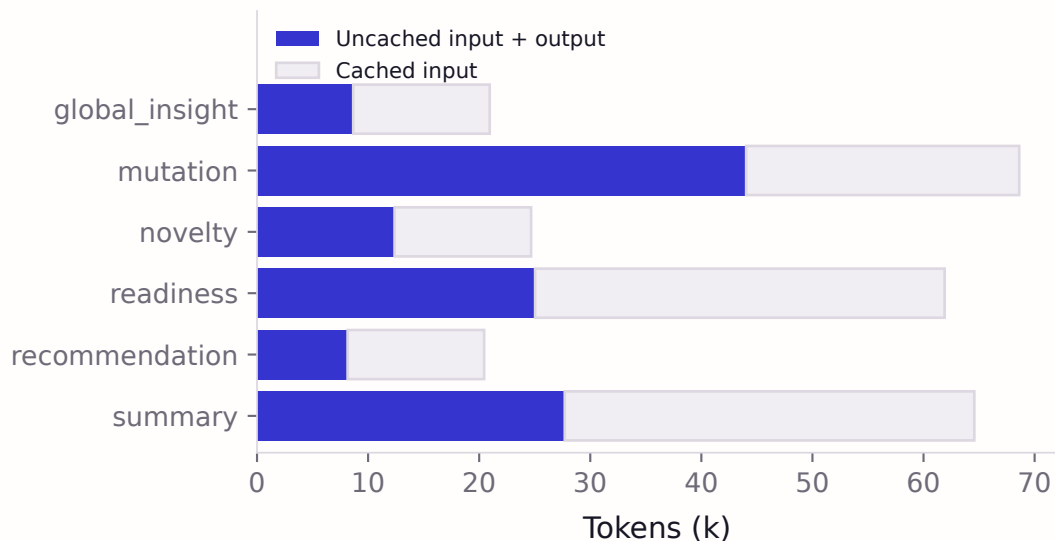
A Progress per reported token use



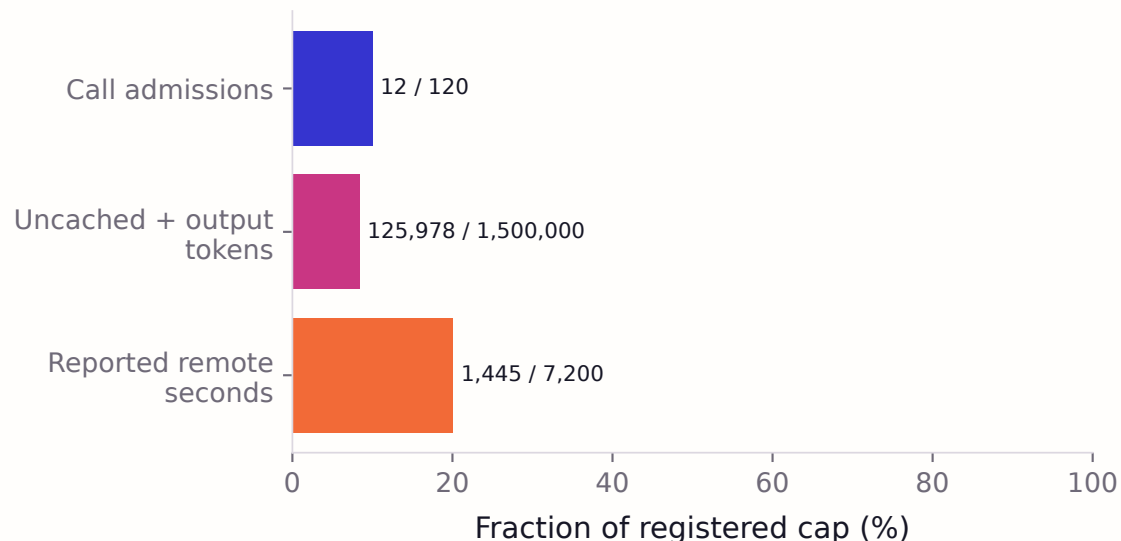
B Measured local episode computation



C Reported experiment-model usage by role



D Shared bounded-execution budget



Readiness/novelty fixture calls consume the shared budget but no candidate slots; the readiness row remains separate from discovery.

1 ended call(s) lack usage: token bars are reported partial sums, not complete cost. Cached input stays separate.

Remote per-call elapsed durations may overlap local work: do not add them to episode CPU or treat them as controller wall time.

Native dollar values are API list-price estimates, not subscription charges. No model-price estimate is plotted as actual spending.