



RUN1 · 3/32 native slots · reused development cases

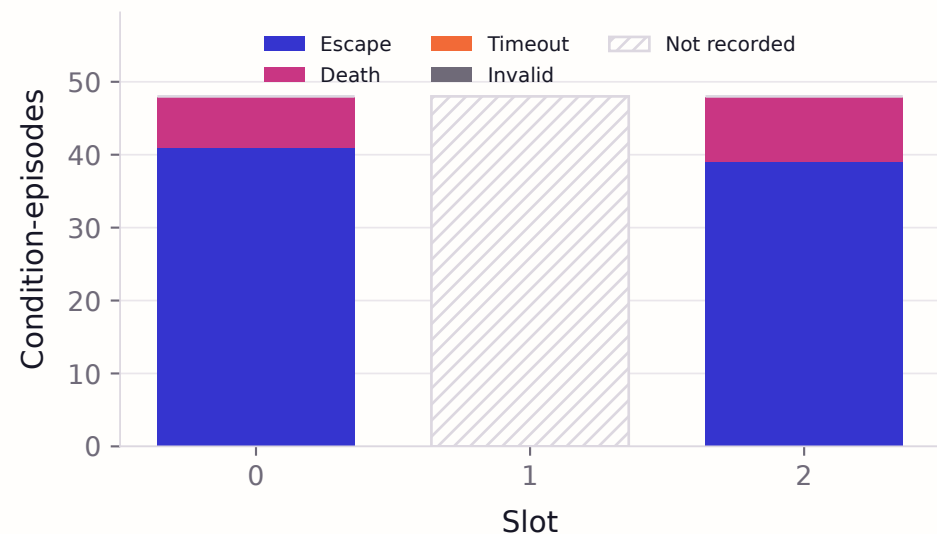
A Every native candidate slot



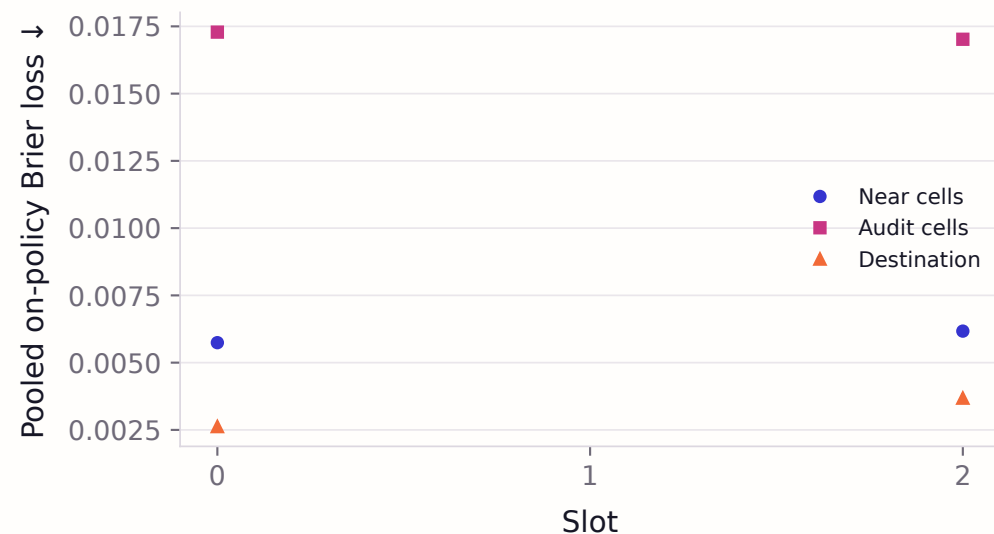
B Task $\neq$ prediction quality



C Outcome denominators stay explicit



D Proper prediction diagnostics



Eight layouts $\times$ six conditions per evaluated slot. Seed copies are administrative rows, not extra candidate evaluations.

Invalid episodes retain zero task; native eligibility additionally requires every episode to validate.

Hatched episodes were not recorded; they are not observed algorithm failures.

Forecasts use different on-policy experience. No matched learning, frozen-control or transfer assessment is established by this figure.